

Problem 1

```
In [ ]: import numpy as np
        from datetime import datetime

        import torch
        import torch.nn as nn
        import torch.nn.functional as F
        from torch.utils.data import DataLoader

        from torchvision import datasets, transforms

        %matplotlib inline
        import matplotlib.pyplot as plt
```

```
In [ ]: # define transforms
        transforms = transforms.Compose([transforms.Resize((32, 32)), transforms.


        # download and create datasets
        train_dataset = datasets.MNIST(
            root="mnist_data", train=True, transform=transforms, download=True
        )

        valid_dataset = datasets.MNIST(root="mnist_data", train=False, transform=
```

1.1.1

```
In [ ]: plt.imshow(train_dataset[0][0].squeeze(), cmap="gray")
        plt.text(10, -2, "The label is " + str(train_dataset[0][1]))
```

```
Out[ ]: Text(10, -2, 'The label is 5')
```



```
In [ ]: # hyper parameters
RANDOM_SEED = 42
LEARNING_RATE = 0.001
BATCH_SIZE = 32
N_EPOCHS = 15

IMG_SIZE = 32
N_CLASSES = 10
```

1.1.2

```
In [ ]: # define the data loaders
train_loader = DataLoader(dataset=train_dataset, batch_size=BATCH_SIZE, s
valid_loader = DataLoader(dataset=valid_dataset, batch_size=BATCH_SIZE, s
```

1.1.3

```
In [ ]: def train(train_loader, model, criterion, optimizer):
    """
    Train one epoch.
    """

    model.train()
    running_loss = 0

    for X, y_true in train_loader:

        optimizer.zero_grad()
```

```

    # Forward pass
    y_hat, _ = model(X)
    loss = criterion(y_hat, y_true)
    running_loss += loss.item() * X.size(0)

    # Backward pass
    loss.backward()
    optimizer.step()

epoch_loss = running_loss / len(train_loader.dataset)
return model, optimizer, epoch_loss

```

1.1.4

```

In [ ]: def validate(valid_loader, model, criterion):
        """
        Function for the validation step of the training loop.
        Returns the model and the loss on the test set.
        """

        model.eval()
        running_loss = 0

        for X, y_true in valid_loader:

            # Forward pass and record loss
            y_hat, _ = model(X)
            loss = criterion(y_hat, y_true)
            running_loss += loss.item() * X.size(0)

        epoch_loss = running_loss / len(valid_loader.dataset)

        return model, epoch_loss

```

```

In [ ]: def training_loop(
        model, criterion, optimizer, train_loader, valid_loader, epochs, print_loss):
        """
        Function defining the entire training loop
        """

        # set objects for storing metrics
        best_loss = 1e10
        train_losses = []
        valid_losses = []
        train_accs = []
        valid_accs = []

        # Train model
        for epoch in range(0, epochs):

            # training
            model, optimizer, train_loss = train(train_loader, model, criterion)
            train_losses.append(train_loss)

            # validation
            with torch.no_grad():
                model, valid_loss = validate(valid_loader, model, criterion)

```

```

        valid_losses.append(valid_loss)

    if epoch % print_every == (print_every - 1):

        train_acc = get_accuracy(
            model,
            train_loader,
        )
        train_accs.append(train_acc)
        valid_acc = get_accuracy(model, valid_loader)
        valid_accs.append(valid_acc)

        print(
            f"{datetime.now().time().replace(microsecond=0)} "
            f"Epoch: {epoch}\t"
            f"Train loss: {train_loss:.4f}\t"
            f"Valid loss: {valid_loss:.4f}\t"
            f"Train accuracy: {100 * train_acc:.2f}\t"
            f"Valid accuracy: {100 * valid_acc:.2f}"
        )

    performance = {
        "train_losses": train_losses,
        "valid_losses": valid_losses,
        "train_acc": train_accs,
        "valid_acc": valid_accs,
    }

    return model, optimizer, performance

```

1.1.5

```

In [ ]: def get_accuracy(model, data_loader):
        """
        Function for computing the accuracy of the predictions over the entire
        dataset.
        """

        correct_pred = 0
        n = 0

        with torch.no_grad():
            model.eval()
            for X, y_true in data_loader:

                _, y_prob = model(X)
                _, predicted_labels = torch.max(y_prob, 1)

                n += y_true.size(0)
                correct_pred += (predicted_labels == y_true).sum()

        return correct_pred.float() / n

def plot_performance(performance):
    """
    Function for plotting training and validation losses
    """

```

```

# temporarily change the style of the plots to seaborn
# plt.style.use("seaborn")

fig, ax = plt.subplots(1, 2, figsize=(16, 4.5))
for key, value in performance.items():
    if "loss" in key:
        ax[0].plot(value, label=key)
    else:
        ax[1].plot(value, label=key)
ax[0].set(title="Loss over epochs", xlabel="Epoch", ylabel="Loss")
ax[1].set(title="accuracy over epochs", xlabel="Epoch", ylabel="Loss")
ax[0].legend()
ax[1].legend()
plt.show()

# change the plot style to default
plt.style.use("default")

```

1.2.1

```

In [ ]: class LeNet5(nn.Module):

    def __init__(self, n_classes):
        super(LeNet5, self).__init__()

        self.feature_extractor = nn.Sequential(
            nn.Conv2d(in_channels=1, out_channels=6, kernel_size=5, stride=1),
            nn.Tanh(),
            nn.AvgPool2d(kernel_size=2),
            nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5, stride=1),
            nn.Tanh(),
            nn.AvgPool2d(kernel_size=2),
            nn.Conv2d(in_channels=16, out_channels=120, kernel_size=5, stride=1),
            nn.Tanh(),
        )
        self.classifier = nn.Sequential(
            nn.Linear(in_features=120, out_features=84),
            nn.Tanh(),
            nn.Linear(in_features=84, out_features=n_classes),
        )

    def forward(self, x):
        x = self.feature_extractor(x)
        x = torch.flatten(x, 1)
        logits = self.classifier(x)
        probs = F.softmax(logits, dim=1)
        return logits, probs

```

1.2.2

```

In [ ]: class MLP(nn.Module):

    def __init__(self, layers):
        super(MLP, self).__init__()

        self.all_layers = nn.ModuleList()

```

```

        for i in range(1, len(layers)):
            self.all_layers.append(
                nn.Linear(in_features=layers[i - 1], out_features=layers[i])
            )
            if i != len(layers) - 1:
                self.all_layers.append(nn.Tanh())
        self.all_layers = nn.Sequential(*self.all_layers)

    def forward(self, x):
        x = x.view(x.shape[0], -1)
        logits = self.all_layers(x)
        probs = F.softmax(logits, dim=1)
        return logits, probs

```

1.3.1

```

In [ ]: torch.manual_seed(RANDOM_SEED)

model = LeNet5(N_CLASSES)
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)
criterion = nn.CrossEntropyLoss()

```

```

In [ ]: model, optimizer, performance_1 = training_loop(
        model, criterion, optimizer, train_loader, valid_loader, N_EPOCHS
    )


```

21:50:18 Epoch: 0	Train loss: 0.2290	Valid loss: 0.1020	Tr
ain accuracy: 96.84	Valid accuracy: 96.81		
21:50:47 Epoch: 1	Train loss: 0.0762	Valid loss: 0.0619	Tr
ain accuracy: 98.39	Valid accuracy: 98.19		
21:51:16 Epoch: 2	Train loss: 0.0550	Valid loss: 0.0542	Tr
ain accuracy: 98.59	Valid accuracy: 98.45		
21:51:45 Epoch: 3	Train loss: 0.0438	Valid loss: 0.0486	Tr
ain accuracy: 99.07	Valid accuracy: 98.41		
21:52:16 Epoch: 4	Train loss: 0.0343	Valid loss: 0.0416	Tr
ain accuracy: 99.24	Valid accuracy: 98.74		
21:52:46 Epoch: 5	Train loss: 0.0297	Valid loss: 0.0440	Tr
ain accuracy: 99.42	Valid accuracy: 98.63		
21:53:15 Epoch: 6	Train loss: 0.0243	Valid loss: 0.0471	Tr
ain accuracy: 99.37	Valid accuracy: 98.51		
21:53:45 Epoch: 7	Train loss: 0.0219	Valid loss: 0.0477	Tr
ain accuracy: 99.43	Valid accuracy: 98.57		
21:54:16 Epoch: 8	Train loss: 0.0191	Valid loss: 0.0402	Tr
ain accuracy: 99.68	Valid accuracy: 98.76		
21:54:45 Epoch: 9	Train loss: 0.0166	Valid loss: 0.0494	Tr
ain accuracy: 99.55	Valid accuracy: 98.68		
21:55:14 Epoch: 10	Train loss: 0.0155	Valid loss: 0.0526	Tr
ain accuracy: 99.42	Valid accuracy: 98.40		
21:55:44 Epoch: 11	Train loss: 0.0135	Valid loss: 0.0418	Tr
ain accuracy: 99.69	Valid accuracy: 98.89		
21:56:13 Epoch: 12	Train loss: 0.0117	Valid loss: 0.0466	Tr
ain accuracy: 99.71	Valid accuracy: 98.77		
21:56:42 Epoch: 13	Train loss: 0.0117	Valid loss: 0.0574	Tr
ain accuracy: 99.65	Valid accuracy: 98.56		
21:57:11 Epoch: 14	Train loss: 0.0122	Valid loss: 0.0504	Tr
ain accuracy: 99.70	Valid accuracy: 98.69		

```

In [ ]: plot_performance(performance_1)

```



1.3.2

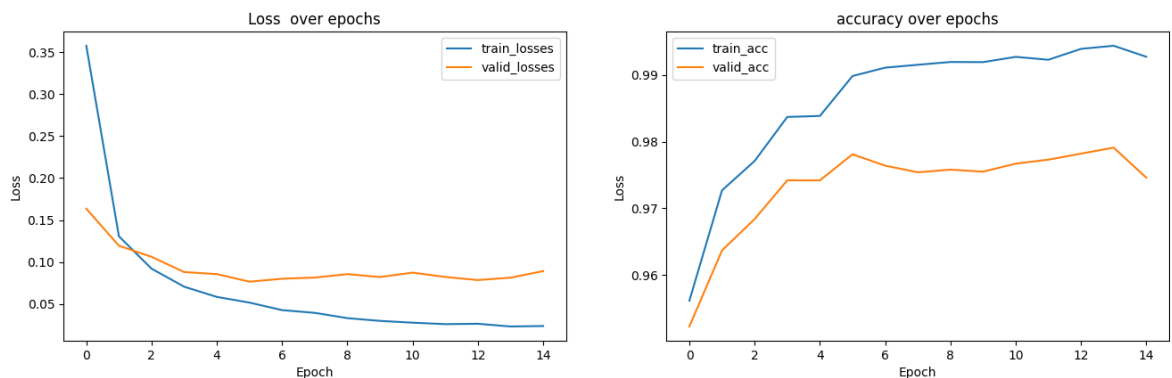
```
In [ ]: torch.manual_seed(RANDOM_SEED)
layers = [1024, 256, 64, 16, N_CLASSES]
model = MLP(layers)
print(model)
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)
criterion = nn.CrossEntropyLoss()
```

```
MLP(
  (all_layers): Sequential(
    (0): Linear(in_features=1024, out_features=256, bias=True)
    (1): Tanh()
    (2): Linear(in_features=256, out_features=64, bias=True)
    (3): Tanh()
    (4): Linear(in_features=64, out_features=16, bias=True)
    (5): Tanh()
    (6): Linear(in_features=16, out_features=10, bias=True)
  )
)
```

```
In [ ]: model, optimizer, performance_2 = training_loop(
    model, criterion, optimizer, train_loader, valid_loader, N_EPOCHS
)
```

21:57:18 Epoch: 0	Train loss: 0.3575	Valid loss: 0.1636	Tr
ain accuracy: 95.61	Valid accuracy: 95.23		
21:57:24 Epoch: 1	Train loss: 0.1307	Valid loss: 0.1195	Tr
ain accuracy: 97.27	Valid accuracy: 96.37		
21:57:31 Epoch: 2	Train loss: 0.0922	Valid loss: 0.1063	Tr
ain accuracy: 97.71	Valid accuracy: 96.84		
21:57:37 Epoch: 3	Train loss: 0.0707	Valid loss: 0.0881	Tr
ain accuracy: 98.37	Valid accuracy: 97.42		
21:57:44 Epoch: 4	Train loss: 0.0586	Valid loss: 0.0858	Tr
ain accuracy: 98.39	Valid accuracy: 97.42		
21:57:50 Epoch: 5	Train loss: 0.0518	Valid loss: 0.0767	Tr
ain accuracy: 98.98	Valid accuracy: 97.81		
21:57:56 Epoch: 6	Train loss: 0.0428	Valid loss: 0.0802	Tr
ain accuracy: 99.11	Valid accuracy: 97.64		
21:58:03 Epoch: 7	Train loss: 0.0396	Valid loss: 0.0816	Tr
ain accuracy: 99.15	Valid accuracy: 97.54		
21:58:09 Epoch: 8	Train loss: 0.0333	Valid loss: 0.0857	Tr
ain accuracy: 99.19	Valid accuracy: 97.58		
21:58:15 Epoch: 9	Train loss: 0.0300	Valid loss: 0.0822	Tr
ain accuracy: 99.19	Valid accuracy: 97.55		
21:58:22 Epoch: 10	Train loss: 0.0279	Valid loss: 0.0874	Tr
ain accuracy: 99.27	Valid accuracy: 97.67		
21:58:28 Epoch: 11	Train loss: 0.0262	Valid loss: 0.0823	Tr
ain accuracy: 99.23	Valid accuracy: 97.73		
21:58:35 Epoch: 12	Train loss: 0.0266	Valid loss: 0.0787	Tr
ain accuracy: 99.39	Valid accuracy: 97.82		
21:58:41 Epoch: 13	Train loss: 0.0234	Valid loss: 0.0815	Tr
ain accuracy: 99.44	Valid accuracy: 97.91		
21:58:47 Epoch: 14	Train loss: 0.0239	Valid loss: 0.0893	Tr
ain accuracy: 99.27	Valid accuracy: 97.46		

```
In [ ]: plot_performance(performance_2)
```



1.4 Comparison of these two models.

1.What is the number of trainable parameters of LeNet? 5 points

- Convolutional Layers 1: $(55+1) * 6 = 156$
- Convolutional Layers 2: $(55*6+1)16 = 2416$
- Convolutional Layers 3: $(55*16+1)120 = 48120$
- Fully Connected Layers 1: $120*84+84 = 10164$
- Fully Connected Layers 2: $84*10+10 = 850$

- Total: 61706

2.What is the number of trainable parameters of MLP? 5 points

- Fully Connected Layers 1: $1024 \times 256 + 256$
- Fully Connected Layers 2: $256 \times 64 + 64$
- Fully Connected Layers 3: $64 \times 16 + 16$
- Last Fully Connected Layers: $16 \times 10 + 10$
- Total: 280058

3.Which model has better performance in terms of prediction accuracy on the test data? Give a reason why this model works better than the other. 10 points

LeNet is better.

When using MLP to process images, it is often necessary to flatten the image from its original form into a one-dimensional vector. This flattening process destroys the spatial relationship between pixels.

LeNet is able to understand the spatial relationship between pixels in an image

Statement of Collaboration (5 points)

I do it by myself.